

# Feedback from surveys

- ◆ Still a little too fast
- ◆ Breakout rooms aren't working
- ◆ What should I know before each lecture?
- ◆ Use Friday's more for review/reinforcement
  - And HW
- ◆ Organize quizzes better on canvas
- ◆ HW2: ....

# **Neural Networks: Deep Learning**

**Lyle Ungar**

**Multilevel network:** architecture, link functions  
**CNNs:** local receptive fields, max pooling

---

**Regularization:**  $L_2$ , early stopping, dropout  
**Gradient Descent** (again)  
**Semi-supervised and transfer learning**

# Quick tensor background

## ◆ What's a tensor?

- As in “tensorflow”
- Or “Tensor Processing Unit” (TPU)
- As in the basic data structure in pytorch
  - Aside: there is a worksheet with more than you need to know about pytorch

# All machine learning is optimization

$$\hat{y} = f(x; \theta)$$

$$\operatorname{argmin}_{\theta} \|y - \hat{y}\|$$

So what's new this decade?

(Slightly) different loss functions

(Slightly) different optimization methods

GPU instead of CPU

Different, flexible, functional forms for  $f$   
which require *regularization*

# Loss functions

$$\hat{y} = f(x; \theta)$$

$$\operatorname{argmin}_{\theta} \|y - \hat{y}\|$$

$$\|y - \hat{y}\|_2$$

$$\|y - \hat{y}\|_1$$

$$\|y - \hat{y}\|_0$$

log-likelihood

cross-entropy (KL-divergence)

later in the course: hinge, exponential

# Flexible model forms

$$\hat{y} = f(\mathbf{x}; \boldsymbol{\theta})$$

$\mathbf{x}$

Web page, ad

Past purchases....

Facebook posts

$y$

Click on ad?

NPV

Age, Sex, Personality, ...

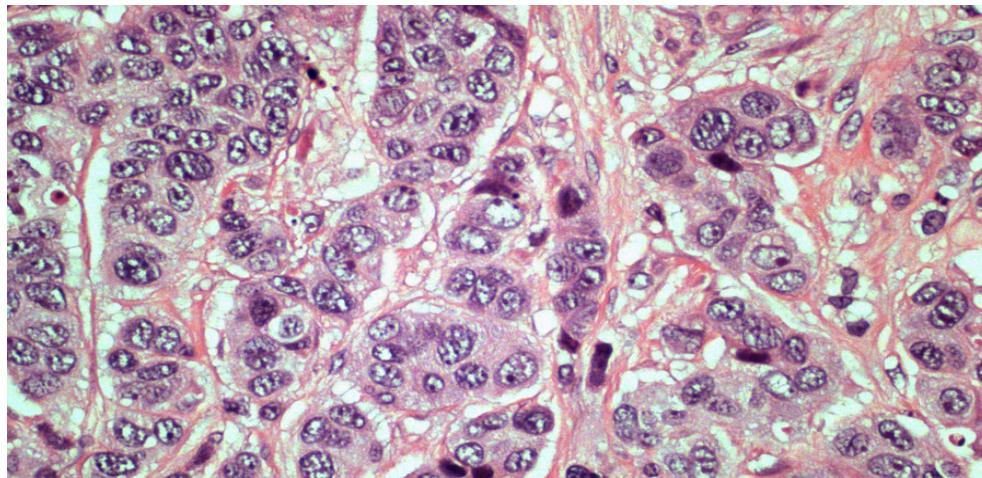
# Flexible model forms

x

biopsy image

y

Cancer present?



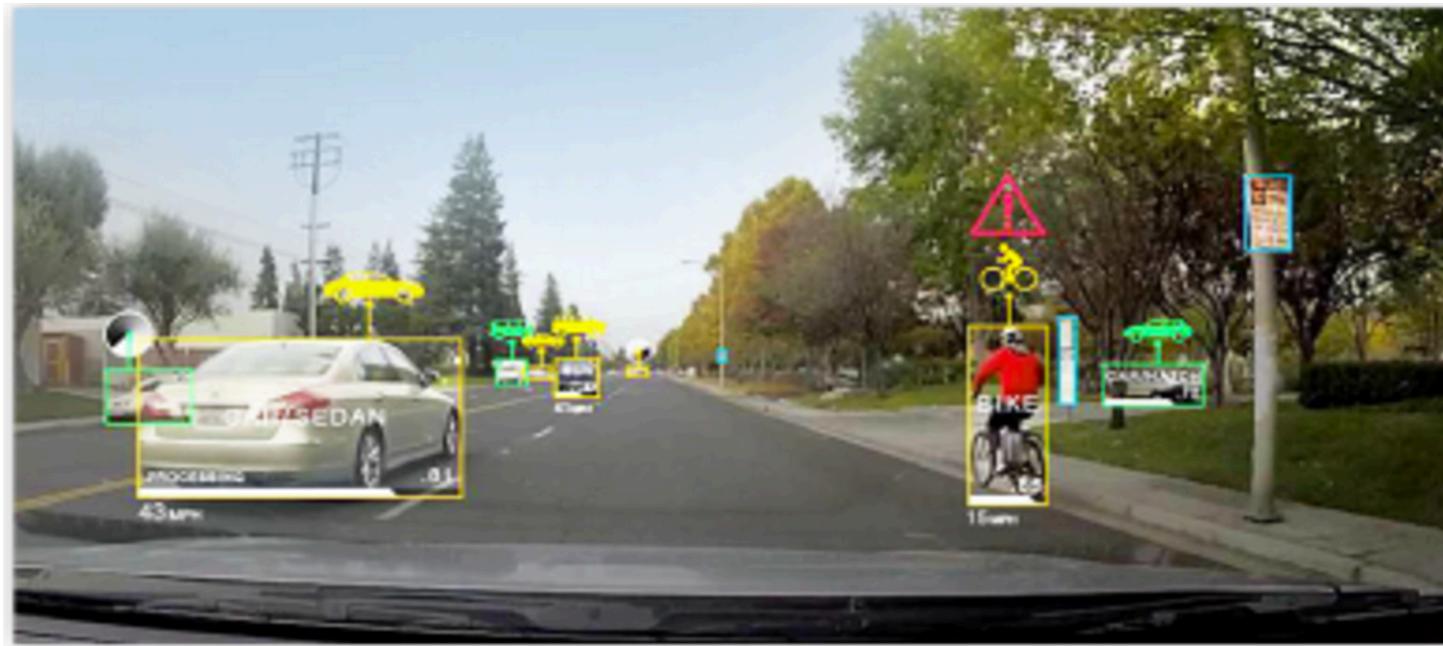
# Flexible model forms

x

Camera image

y

Objects in it



nvidia

# Flexible model forms

x

English sentence

y

Translation

English - detected ▾



Arabic ▾



I love machine  
learning

Edit

أَحُبُّ تَعْلِمُ الْآلاتَ

'uhibb taelam alala

[Open in Google Translate](#)

*Feedback*

# Artificial Neural Nets

- ◆ **Semi-parametric**

- Flexible model form

- ◆ **Used when there are vast amounts of data**

- Hence popular (again) now
  - But recently with smaller training sets.

- ◆ **Deep networks**

- Idea: representation should have *many* different levels of abstraction

# Neural Nets can be

- ◆ **Supervised**

- Generalizes *logistic regression* to a semi-parametric form

- ◆ **Unsupervised**

- Generalizes *PCA* to a semi-parametric form

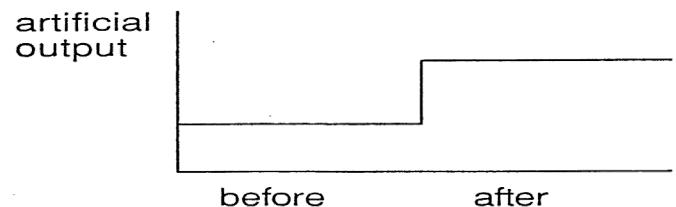
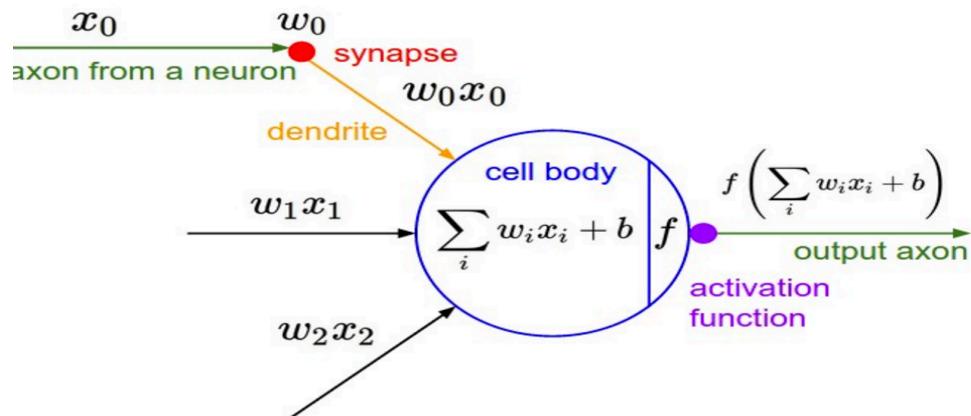
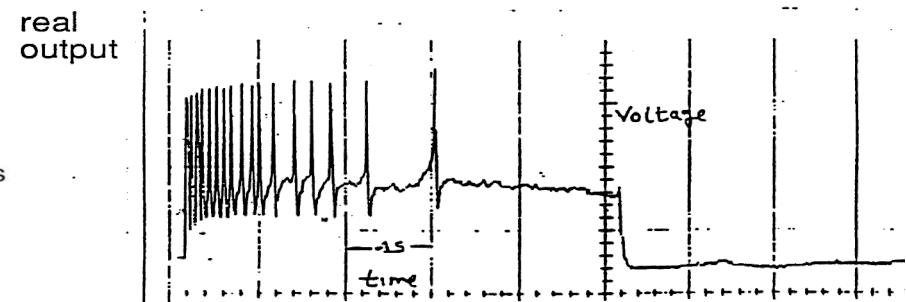
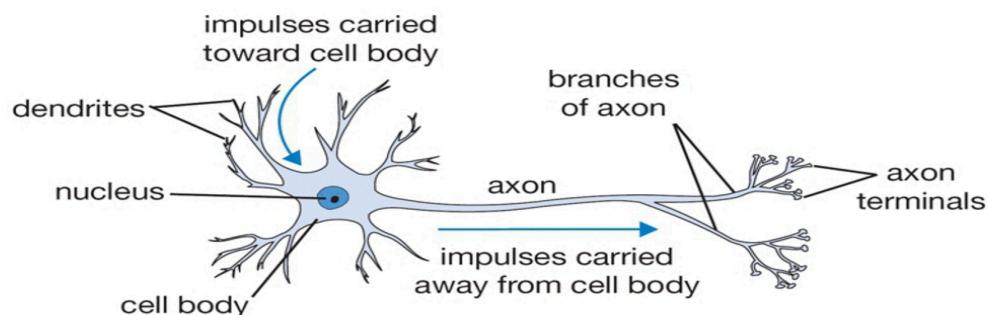
- ◆ **Adversarial**

- ◆ **Semi-supervised**

- ◆ **Reinforcement**

**Neural nets often have built-in structure**

# “Real” and Artificial neuron

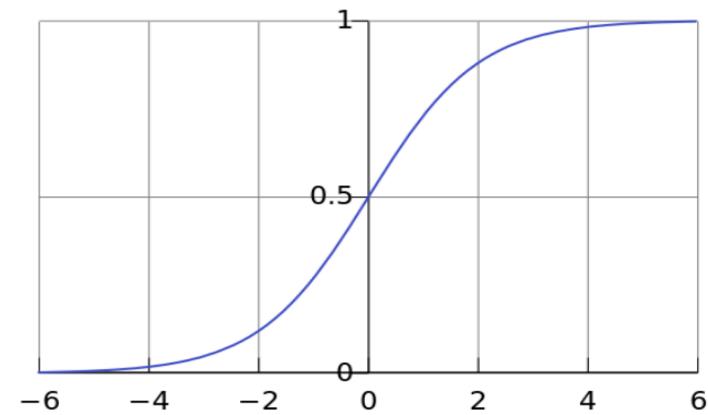
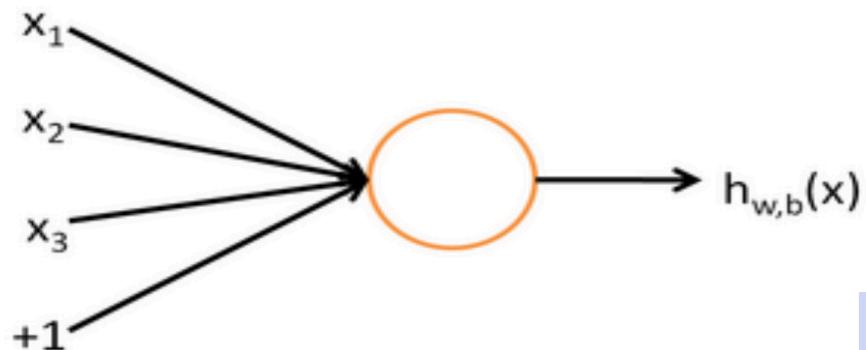


<http://cs231n.github.io/neural-networks-1/>

# One neuron does logistic regression

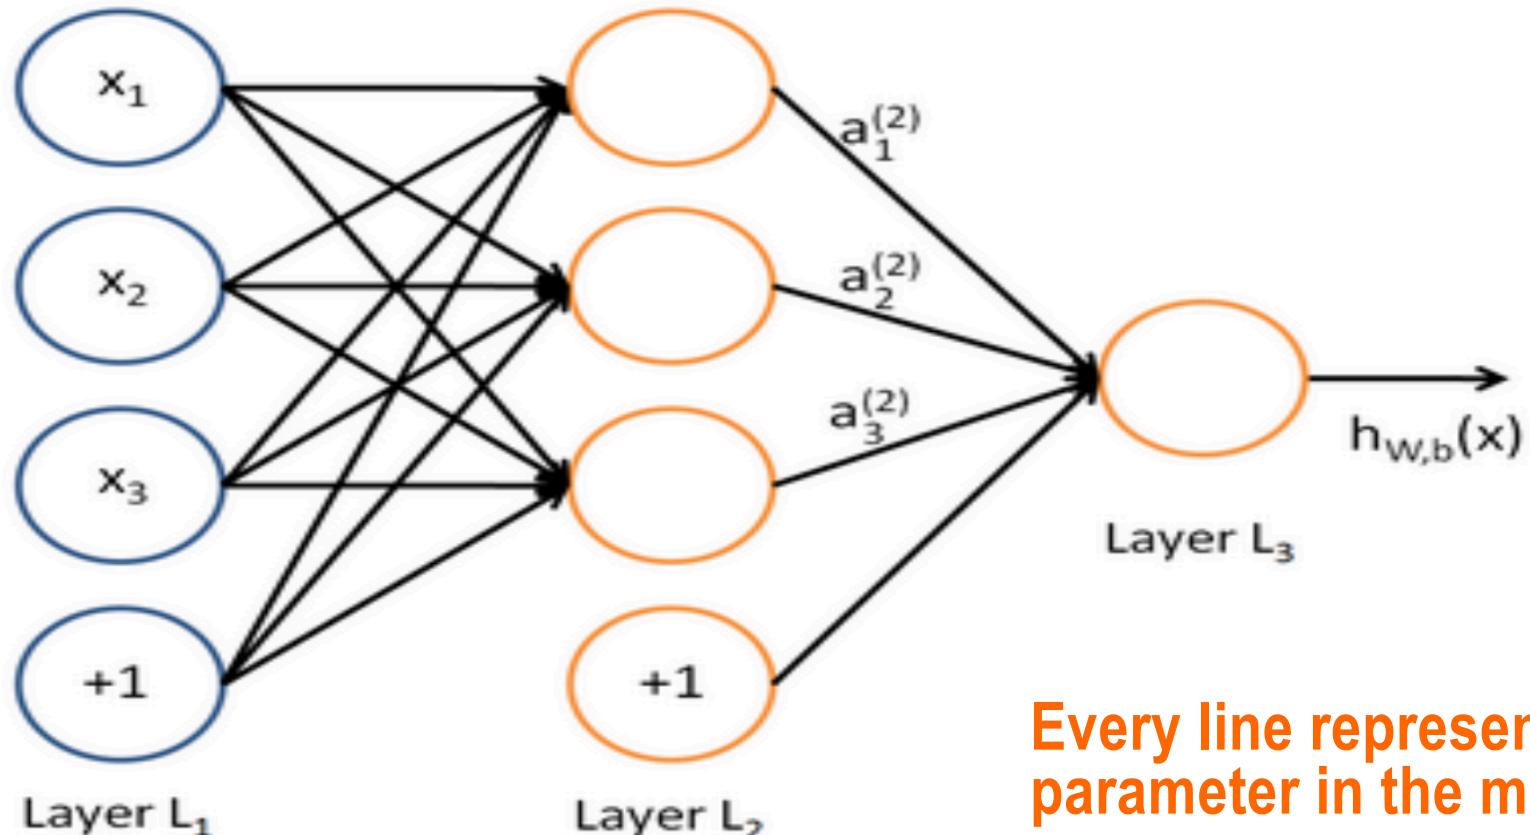
$$h_{w,b}(x) = f(w^\top x + b)$$

$$f(z) = \frac{1}{1 + e^{-z}}$$



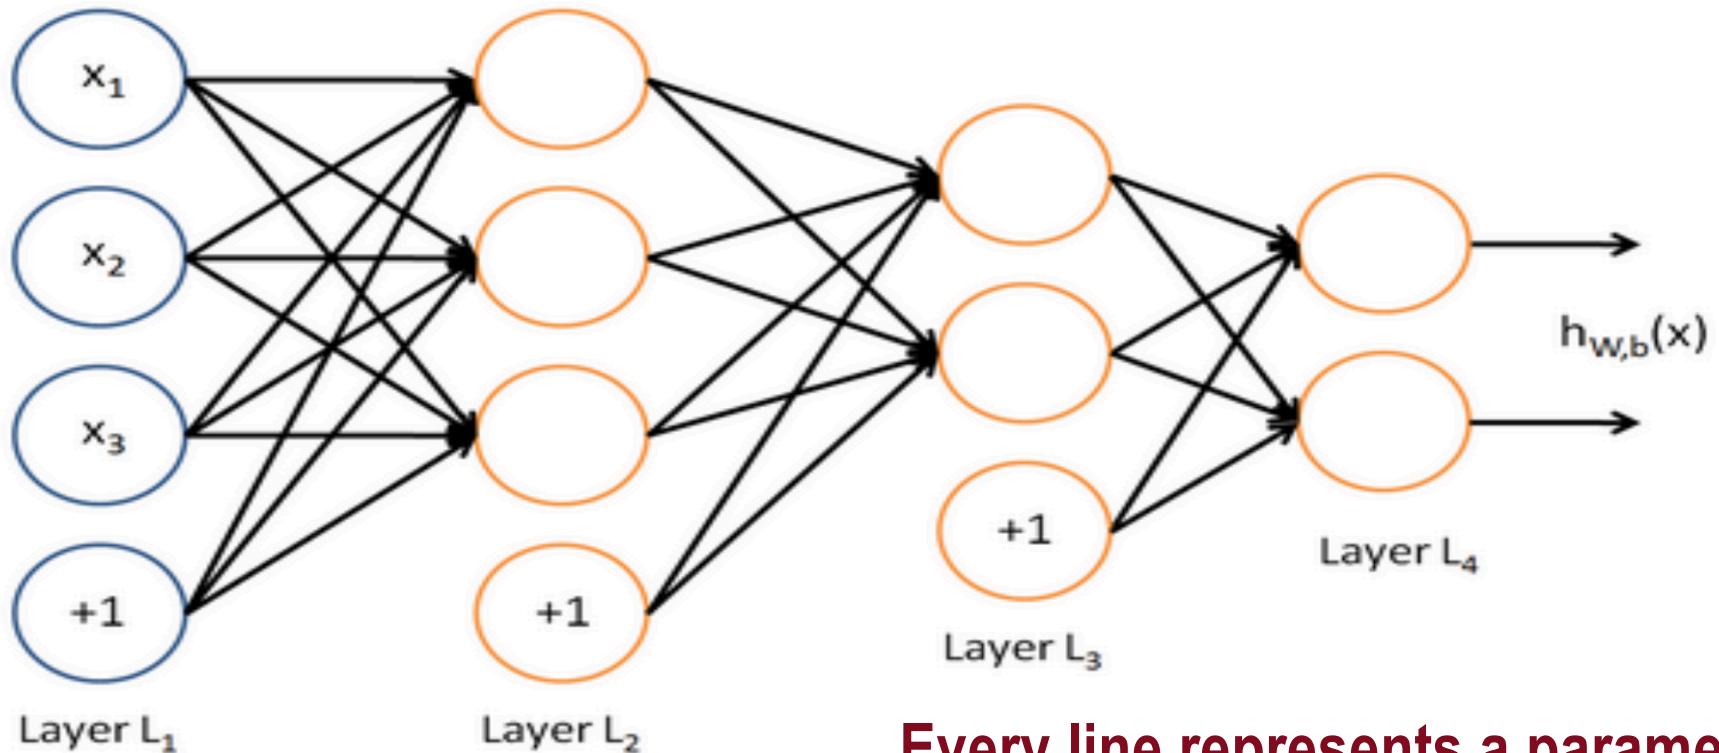
Socher and Manning tutorial

# Neural nets stack logistic regressions



Every line represents a parameter in the model

# Neural nets stack logistic regressions



**Every line represents a parameter  
in the model**

# Training

- ◆ Mini-batch gradient descent
- ◆ “Backpropagate” error derivatives through the model = chain rule

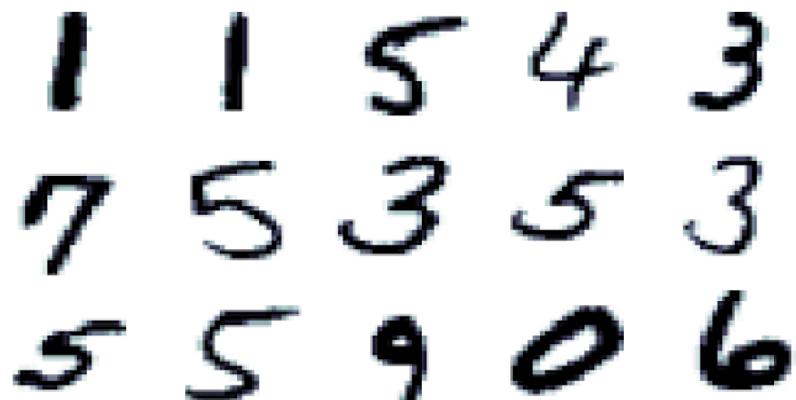
# ANNs do pattern recognition

## ◆ Map input “percepts” to output categories or actions

- Image of an object → what it is
- Image of a person → who it is
- Picture → caption describing it
- Board position → probability of winning
- A word → the sound of saying it
- Sound of a word → the word
- Sequence of words in English → their Chinese translation

# MNIST

- Classify 28x28 images of handwritten digits
- **Train:** 50,000
- **Test:** 10,000



Error (%)	Method	Reference
5.0	KNN	Lecun et al. (1998)
3.6	1k RBF + linear classifier	Lecun et al. (1998)
1.6	2-layer NN	Simard et al. (2003)
1.53	boosted stumps	Kegl et al. (2009)
1.4	SVM	Lecun et al. (1998)
0.79	DNN	Srivastava (2013)
0.45	conv-DNN	Goodfellow et al. (2013)
0.21	conv-DNN	Wi et al. (2013)

# Street View House Numbers

- Classify 32x32 color images of digits
- Digits taken from housenumbers in Google Street View
- **Train:** 604,388
- **Test:** 26,032

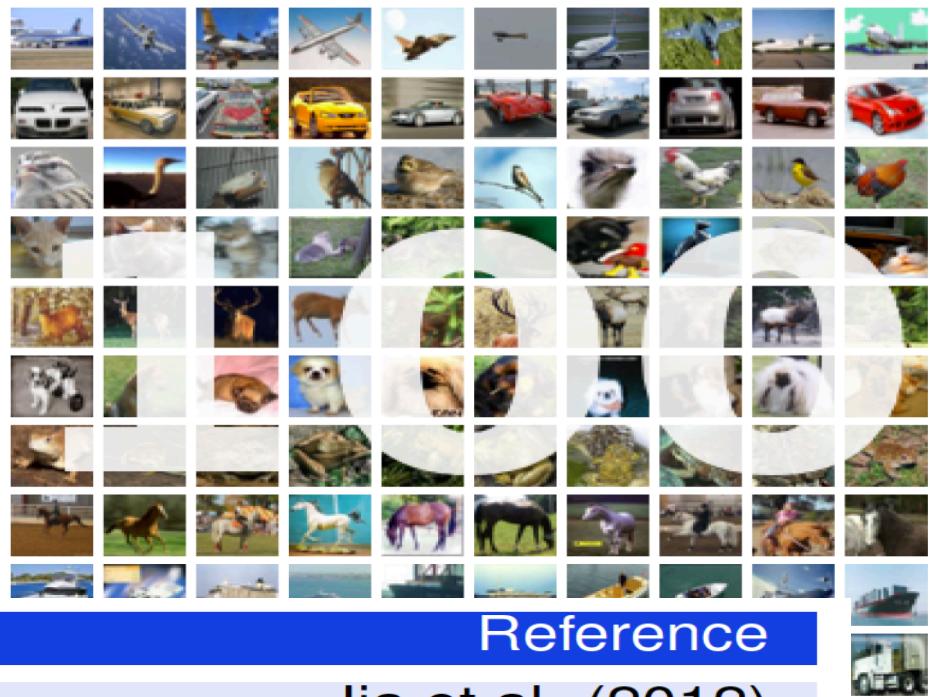


Error (%)	Method	Reference
36.7	WDCH	Netzer et al. (2011)
15	HOG	Netzer et al. (2011)
9.4	KNN	Netzer et al. (2011)
2.47	conv-DNN	Goodfellow et al. (2013)
2	Human	Netzer et al. (2013)
1.92	conv-DNN	Lee et al. (2015)



# CIFAR-100

- Classify 32x32 color images into 100 classes
- Images taken from TinyImages dataset at MIT
- **Train:** 50,000
- **Test:** 10,000



Error (%)	Method	Reference
43.77	SVM	Jia et al. (2012)
39.20	OMP	Lin and Kung (2014)
38.57	conv-DNN	Goodfellow et al. (2013)
36.18	DNN	Srivastava and Alakhutdinov (2015)
34.57	conv-DNN	Lee et al. (2015)

# ImageNet Classification with Deep Convolutional Neural Networks

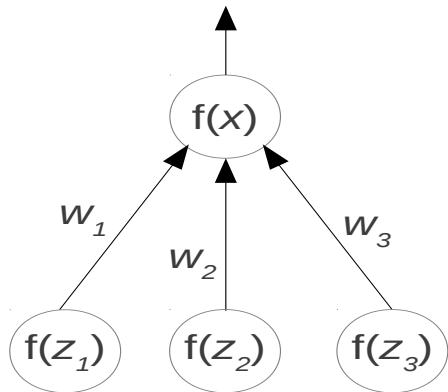
Alex Krizhevsky  
Ilya Sutskever  
Geoffrey Hinton

University of Toronto  
Canada

“AlexNet”      2012

# Neural networks

- A neuron



$$x = w_1 f(z_1) + w_2 f(z_2) + w_3 f(z_3)$$

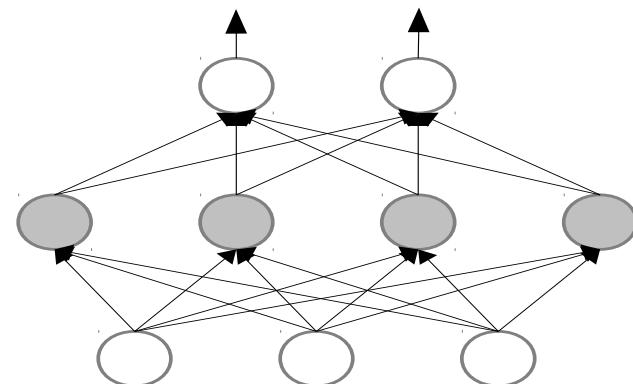
$x$  is called the total input to the neuron, and  $f(x)$  is its output

- A neural network

Output

Hidden

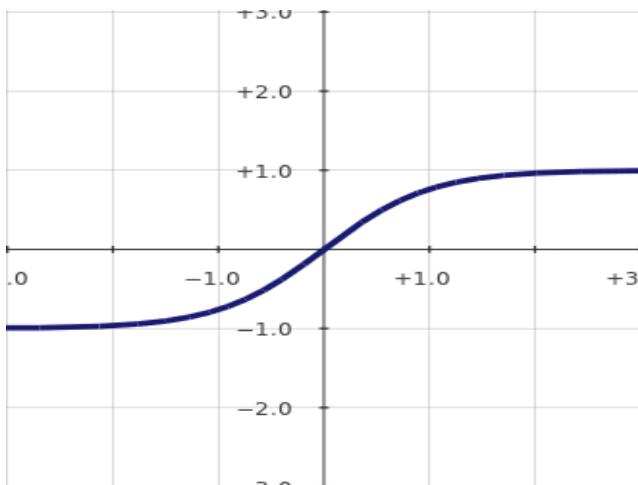
Data



A neural network computes a differentiable function of its input. For example, ours computes:  
 $p(\text{label} \mid \text{an input image})$

**Traditional: sigmoidal**  
e.g. logistic function

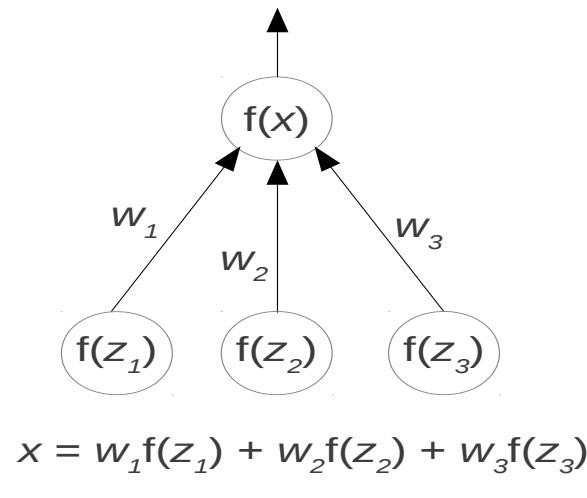
$$f(x) = \tanh(x)$$



**Hyperbolic tangent**

Very bad (slow to train)

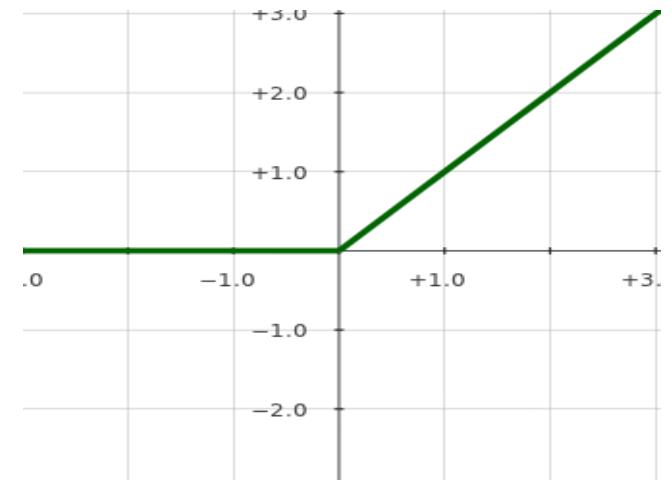
# Neurons



$x$  is called the total input  
to the neuron, and  $f(x)$   
is its output

**But one can use any  
nonlinear function**

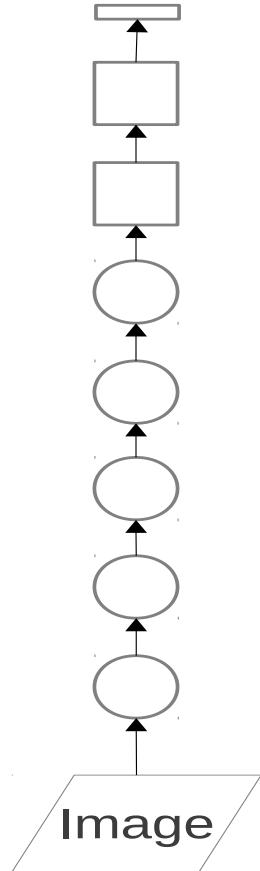
$$f(x) = \max(0, x)$$



**Rectified Linear Unit (ReLU)**

Very good (quick to train)

# Overview of our model



- **Deep:** 7 hidden “weight” layers
- **Learned:** all feature extractors initialized at white Gaussian noise and learned from the data
- Entirely supervised
- **More data = good**



**Convolutional layer:** convolves its input with a bank of 3D filters, then applies point-wise non-linearity



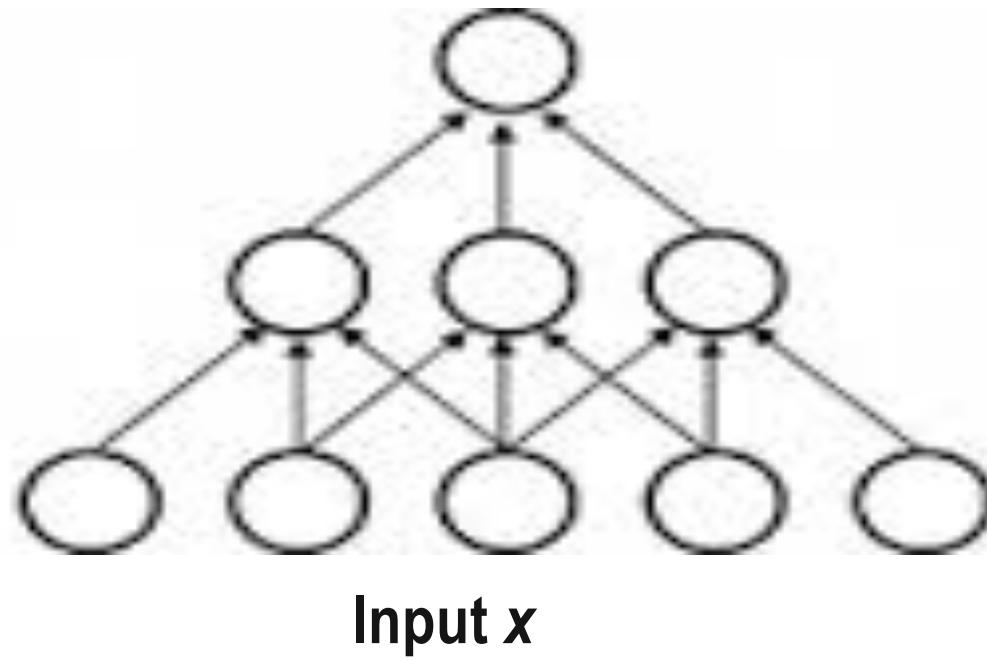
**Fully-connected layer:** applies linear filters to its input, then applies point-wise non-linearity

# Local receptive fields

layer  $m+1$

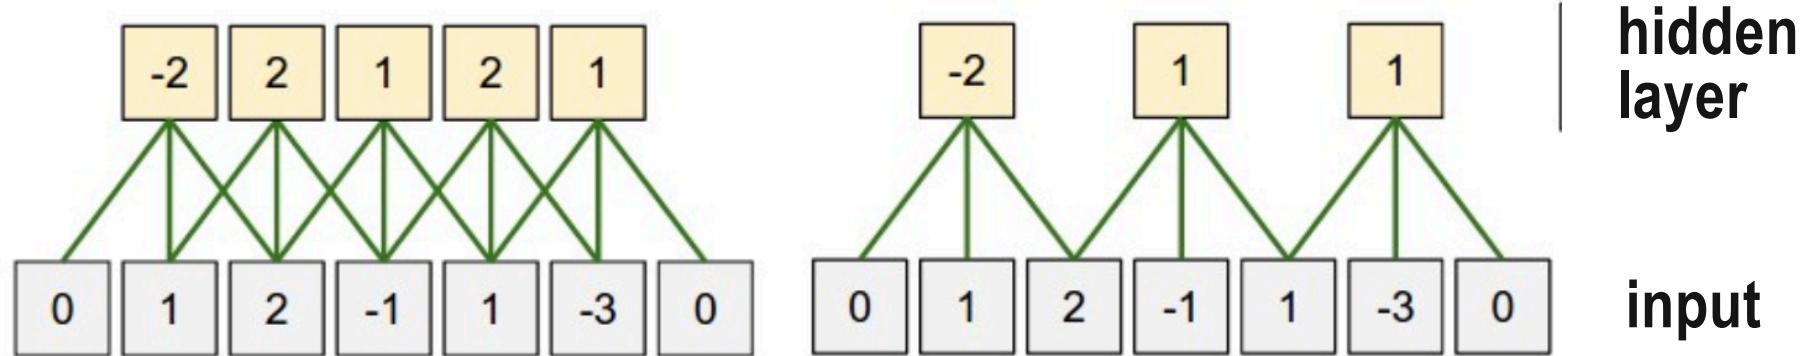
layer  $m$

layer  $m-1$



In vision, a neuron may only get inputs from a limited set of “nearby” neurons

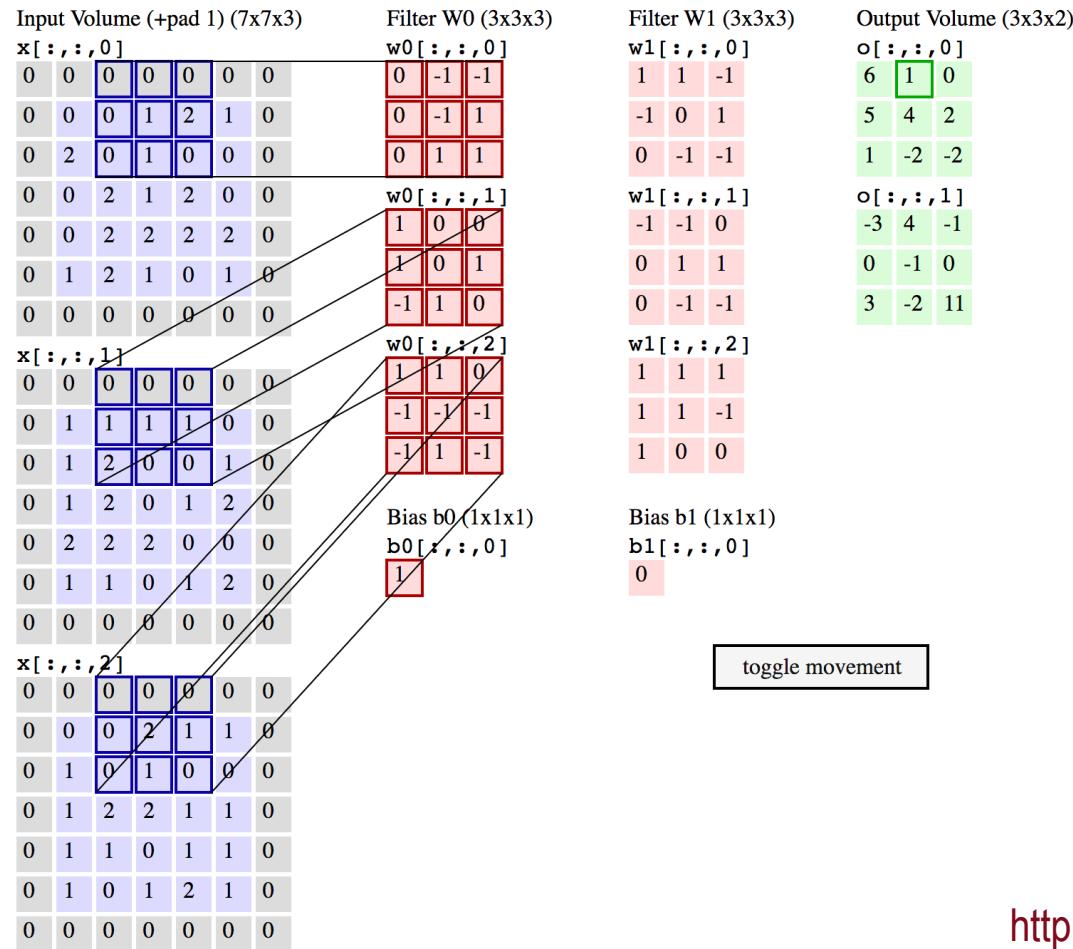
# Local receptive fields



- spatial extent  $F = 3$
- stride  $S = 1$

- spatial extent  $F = 3$
- stride  $S = 2$

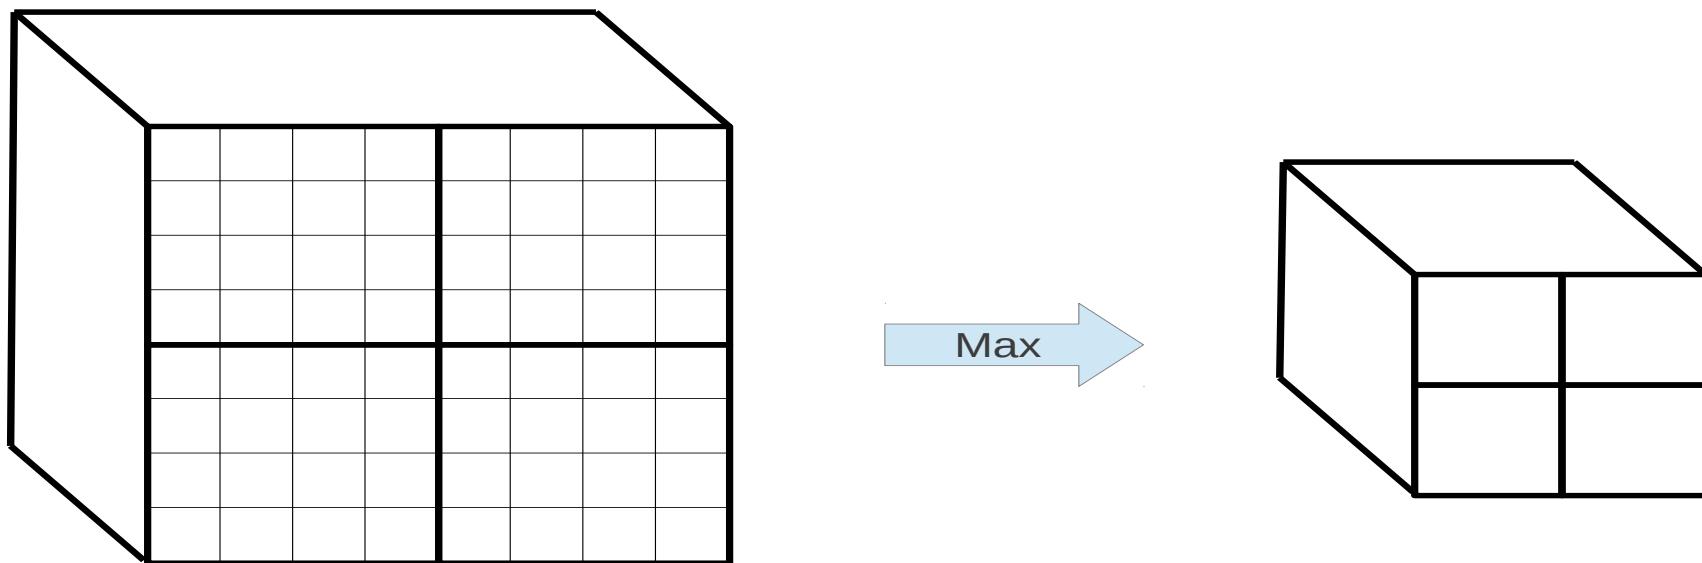
# Local receptive fields



<http://cs231n.github.io/convolutional-networks/>

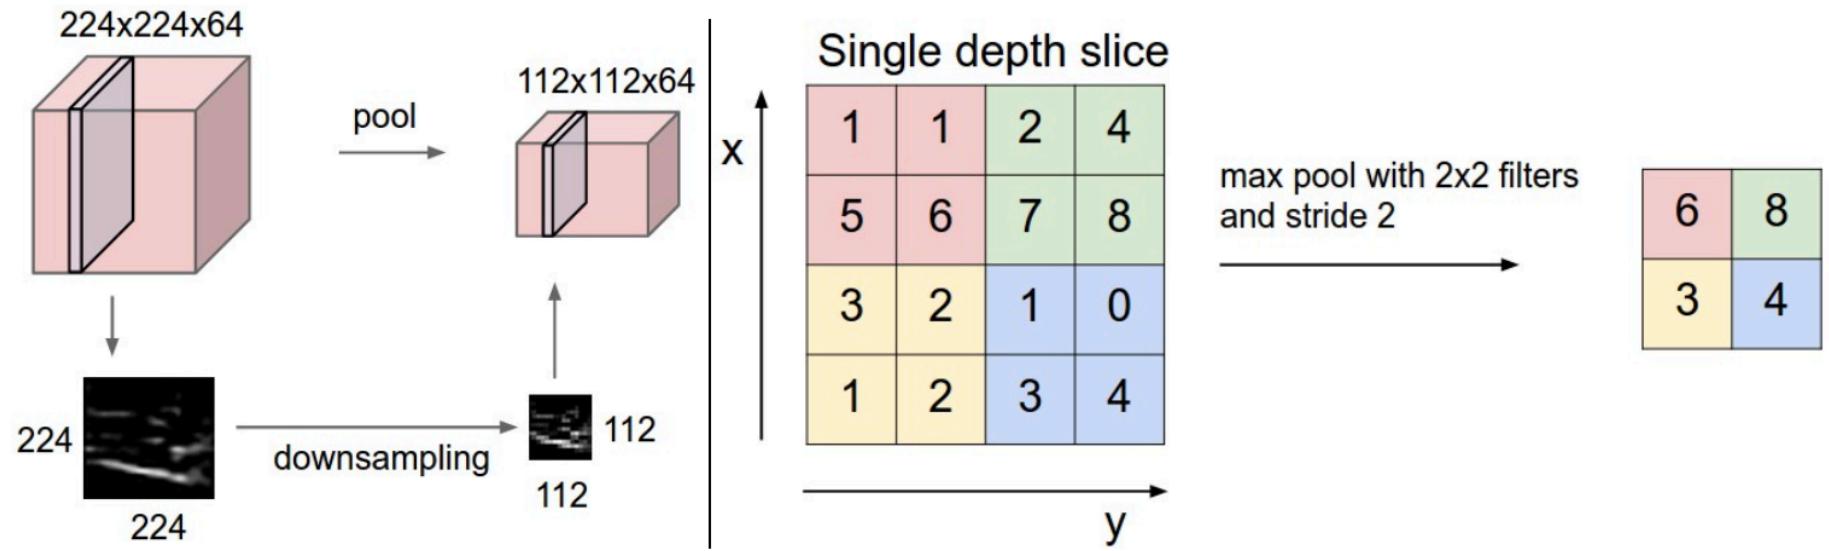
# Local pooling

**Max-pooling** partitions the input image into local regions and outputs the maximum value for each.



**Reduces the computational complexity**  
**Provides translation invariance.**

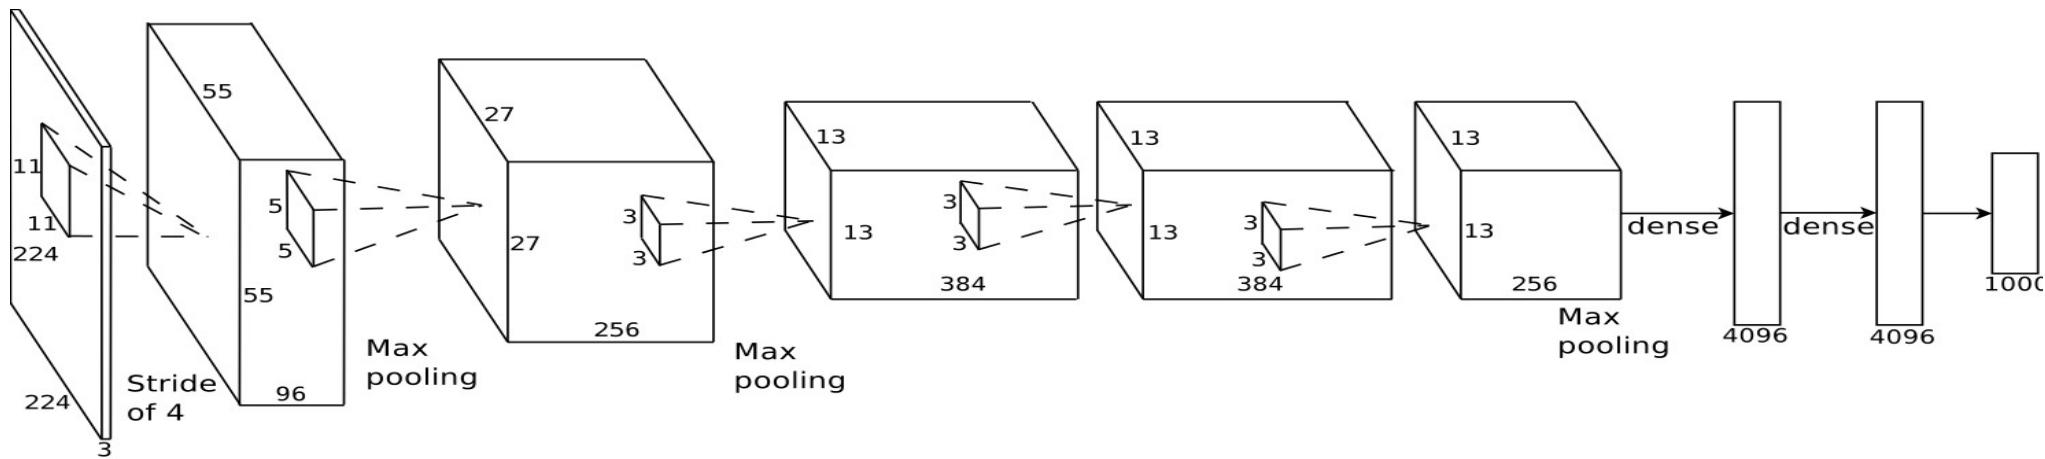
# Max pooling



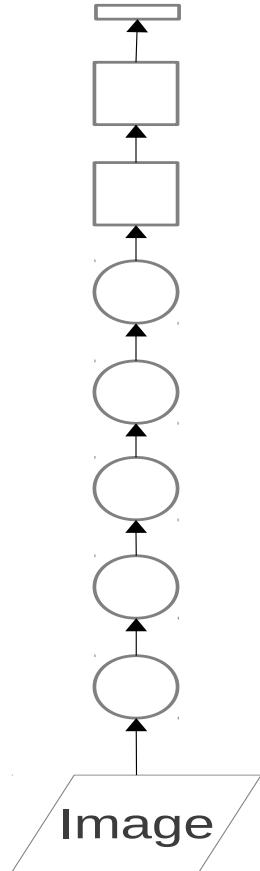
- *Spatial extent  $F=2$*
- *Stride  $S=2$*

# Our model

- Max-pooling layers follow first, second, and fifth convolutional layers
- The number of neurons in each layer is given by 253440, 186624, 64896, 64896, 43264, 4096, 4096, 1000



# Overview of our model



- Trained with stochastic gradient descent on two NVIDIA GPUs for about a week
- 650,000 neurons
- 60,000,000 parameters
- 630,000,000 connections
- **Final feature layer:** 4096-dimensional

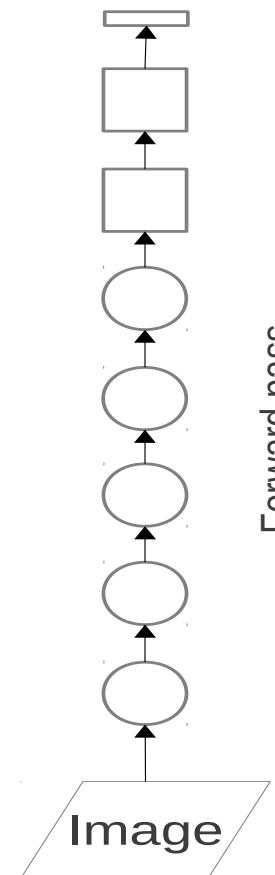
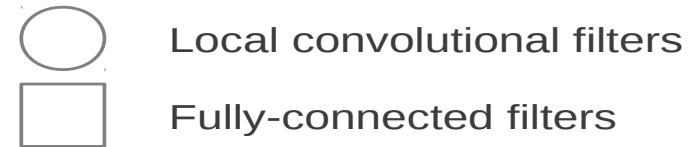


**Convolutional layer:** convolves its input with a bank of 3D filters, then applies point-wise non-linearity



**Fully-connected layer:** applies linear filters to its input, then applies point-wise non-linearity

# Training



Using stochastic gradient descent and the *backpropagation algorithm* (just repeated application of the chain rule)

One output unit per class

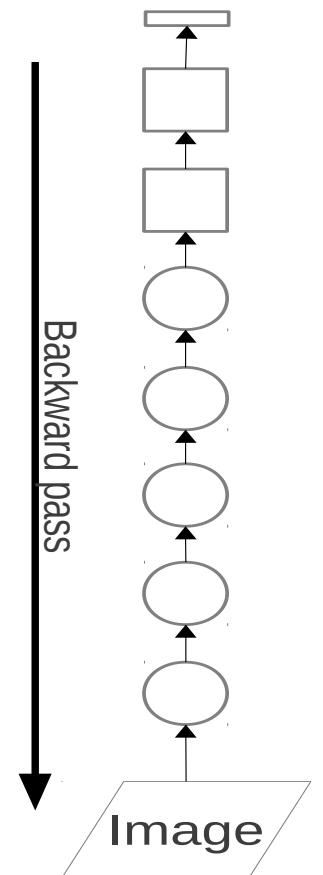
$x_i$  = total input to output unit  $i$

$$f(x_i) = \frac{\exp(x_i)}{\sum_{j=1}^{1000} \exp(x_j)}$$

We maximize the log-probability of the correct label,  $\log f(x_t)$

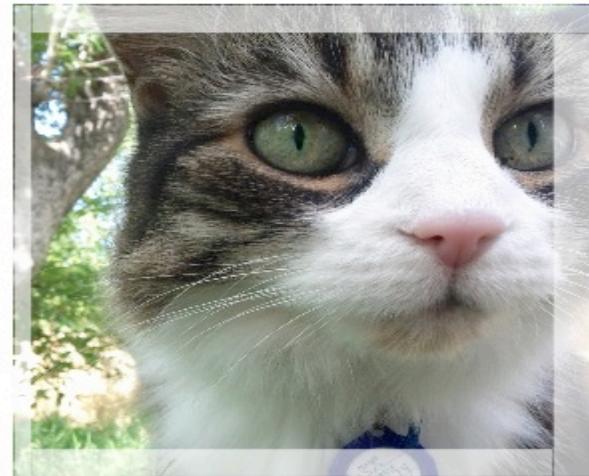
$f(x_i)$  = softmax

$\log(f(x_t))$  = cross-entropy



# Data augmentation

- Our neural net has 60M real-valued parameters and 650,000 neurons
- It overfits a lot. Therefore we train on 224x224 patches extracted randomly from 256x256 images, and also their horizontal reflections.





# Questions?

Top



Start the presentation to see live content. For screen share software, share the entire screen. Get help at [pollev.com/app](https://pollev.com/app)