

# Online Learning: LMS and Perceptrons

Partially adapted from slides by Ryan Gabbard  
and Mitch Marcus (and lots original slides by Lyle Ungar)

## Learning Objectives

Complexity of OLS

LMS = SDG

Perceptron variations


online hinge loss optimization

Note: not on midterm

# Why do online learning?

- ◆ Batch learning can be expensive for big datasets
  - How expensive is it to compute  $(X^T X)^{-1}$  ?

A)  $n^3$   
B)  $p^3$   
C)  $np^2$   
D)  $n^2p$



# Why do online learning?

## ◆ Batch learning can be expensive for big datasets

- How hard is it to compute  $(X^T X)^{-1}$ ?
  - $np^2$  to form  $X^T X$
  - $p^3$  to invert
- Tricky to parallelize inversion

## ◆ Online methods are easy in a map-reduce environment

- They are often clever versions of stochastic gradient descent

Have you seen map-reduce/hadoop?

- A) Yes
- B) No



# Online learning methods

## ◆ Least Mean Squares (LMS)

- Online regression --  $L_2$  error

## ◆ Perceptron

- Online SVM -- Hinge loss

# LMS: Online linear regression

- ◆ Minimize  $\text{Err} = \sum_i (y_i - \mathbf{w}^T \mathbf{x}_i)^2$  using stochastic gradient descent

- Look at each observation  $(\mathbf{x}_i, y_i)$  sequentially and decrease its error  $\text{Err}_i = (y_i - \mathbf{w}^T \mathbf{x}_i)^2$

- ◆ LMS (Least Mean Squares) algorithm

- $\mathbf{w}_{i+1} = \mathbf{w}_i - \eta/2 \frac{d\text{Err}_i}{d\mathbf{w}_i}$
  - $d\text{Err}_i/d\mathbf{w}_i = -2 (y_i - \mathbf{w}_i^T \mathbf{x}_i) \mathbf{x}_i = -2 r_i \mathbf{x}_i$

$$\mathbf{w}_{i+1} = \mathbf{w}_i + \eta r_i \mathbf{x}_i$$

How do you pick the “learning rate”  $\eta$ ?

Note that  $i$  is the index for both the iteration and the observation, since there is one update per observation

# Online linear regression

- ◆ LMS (Least Mean Squares) algorithm

$$\mathbf{w}_{i+1} = \mathbf{w}_i + \eta \mathbf{r}_i \mathbf{x}_i$$

- ◆ Converges for  $0 < \eta < \lambda_{\max}$ 
  - Where  $\lambda_{\max}$  is the largest eigenvalue of the covariance matrix  $\mathbf{X}^T \mathbf{X}$
- ◆ Convergence rate is proportional to  $\lambda_{\min}/\lambda_{\max}$   
(ratio of extreme eigenvalues of  $\mathbf{X}^T \mathbf{X}$ )

# Perceptron Learning Algorithm

**Input:** A list  $T$  of training examples  $\langle \vec{x}_0, y_0 \rangle \dots \langle \vec{x}_n, y_n \rangle$  where  
 $\forall i : y_i \in \{+1, -1\}$

**Output:** A classifying hyperplane  $\vec{w}$

Randomly initialize  $\vec{w}$ ;

**while** model  $\vec{w}$  makes errors on the training data **do**

**for**  $\langle \vec{x}_i, y_i \rangle$  in  $T$  **do**

        Let  $\hat{y} = \text{sign}(\vec{w} \cdot \vec{x}_i)$ ;

**if**  $\hat{y} \neq y_i$  **then**

$\vec{w} = \vec{w} + y_i \vec{x}_i$ ;

If you were wrong, make  
 $w$  look more like  $x$

**end**

**end**

**end**

What do we do if error is zero?

Of course, this only converges for linearly separable data

# Perceptron Learning Algorithm

For each observation  $(\mathbf{x}_i, y_i)$

$$\mathbf{w}_{i+1} = \mathbf{w}_i + \eta r_i \mathbf{x}_i$$

Where  $r_i = y_i - \text{sign}(\mathbf{w}_i^T \mathbf{x}_i)$


and  $\eta = 1/2$

I.e., if we get it right: *no change*

if we got it wrong:  $\mathbf{w}_{i+1} = \mathbf{w}_i + y_i \mathbf{x}_i$

**How does this relate to SVMs?**


# Perceptron update



If the prediction at  $\vec{x}_1$  is wrong,  
what is the true label  $y_1$ ?

How do you update  $\vec{w}$ ?

# Perceptron update example



# Properties of the simple perceptron

## ◆ Provably:

- If it's possible to separate the data with a hyperplane (i.e. if it's **linearly separable**), then the algorithm will converge to that hyperplane.
- And it will converge such that the number of mistakes  $M$  it makes is bounded by

$$M < R^2/\gamma^2$$

where

$$R = \max_i \|\mathbf{x}_i\|_2 \quad \text{size of biggest } \mathbf{x}$$

$$\gamma > y_i \mathbf{w}^{*T} \mathbf{x}_i \quad > 0 \text{ if separable}$$

# Properties of the Simple Perceptron

**But what if it isn't separable?**

- Then perceptron is unstable and bounces around

# Voted Perceptron

- ◆ Works just like a regular perceptron, except you keep track of all the intermediate models you created
- ◆ When you want to classify something, you let each of the many models *vote* on the answer and take the *majority*

Often implemented after a “burn-in” period

# Properties of Voted Perceptron

- ◆ **Simple!**
- ◆ **Much better generalization performance than regular perceptron**
  - Almost as good as SVMs
  - Can use the ‘kernel trick’ – replace dot product with another kernel
- ◆ **Training is as fast as a regular perceptron**
- ◆ **But run-time is slower**
  - Since we need  $n$  models

# Averaged Perceptron

- ◆ The final model is the *average* of all the intermediate models
- ◆ Approximation to voted perceptron
- ◆ Again extremely simple!
  - and can use kernels
- ◆ Nearly as fast to train and exactly as fast to run as regular perceptron

# Many possible perceptrons

- ◆ If point  $x_i$  is misclassified
  - $w_{i+1} = w_i + \eta y_i x_i$
- ◆ Different ways of picking *learning rate*  $\eta$
- ◆ Standard perceptron:  $\eta = 1$ 
  - Guaranteed to converge to the correct answer in a finite time if the points are separable (but oscillates otherwise)
  - Can get bounds on error even for non-separable case
- ◆ Alternate: pick  $\eta$  to maximize the margin ( $w_i^T x_i$ ) in some fashion

# Can we do a better job of picking $\eta$ ?

## ◆ Perceptron:

For each observation  $(y_i, \mathbf{x}_i)$

$$\mathbf{w}_{i+1} = \mathbf{w}_i + \eta r_i \mathbf{x}_i$$

$$\text{where } r_i = y_i - \text{sign}(\mathbf{w}_i^T \mathbf{x}_i)$$

$$\text{and } \eta = 1/2$$

Let's use the fact that we are actually trying to minimize a loss function

# Passive Aggressive Perceptron

- Minimize the *hinge loss* at each observation
  - $L(\mathbf{w}_i; \mathbf{x}_i, y_i) = 0$  if  $y_i \mathbf{w}_i^T \mathbf{x}_i \geq 1$  (loss 0 if correct with margin  $> 1$ )  
 $1 - y_i \mathbf{w}_i^T \mathbf{x}_i$  else
- Pick  $\mathbf{w}_{i+1}$  to be as close as possible to  $\mathbf{w}_i$  while still setting the hinge loss to zero
  - If point  $\mathbf{x}_i$  is correctly classified with a margin of at least 1
    - no change
  - Otherwise
    - $\mathbf{w}_{i+1} = \mathbf{w}_i + \eta y_i \mathbf{x}_i$
    - where  $\eta = L(\mathbf{w}_i; \mathbf{x}_i, y_i) / \|\mathbf{x}_i\|^2$
- Can prove bounds on the total hinge loss


# Passive-Aggressive = MIRA

$$w_{i+1} = w_i + \frac{y_i - w_i \cdot x_i}{\|x_i\|^2} x_i$$

easy to show:

$$y_i(w_{i+1} \cdot x_i) = y_i \left( w_i + \frac{y_i - w_i \cdot x_i}{\|x_i\|^2} x_i \right) \cdot x_i = 1$$

new score       $y_i (w_i \cdot x_i + y_i - w_i \cdot x_i) = y_i y_i$



Moves hyperplane so that new point is on the margin

# Margin-Infused Relaxed Algorithm (MIRA)

- ◆ **Multiclass**; each class has a prototype vector
  - Note that the prototype  $w$  is like a feature vector  $x$
- ◆ Classify an instance by choosing the class whose prototype vector is *most similar* to the instance
  - *Has the greatest dot product with the instance*
- ◆ During training, make the ‘smallest’ change to the prototype vectors which guarantees correct classification by a specified margin
  - “passive aggressive”

# Can we parallelize SGD?

- ◆ If I give you 1,000 machines, how do you speed SGD up?

# What we didn't cover

- ◆ Feature selection

# What you should know

## ◆ LMS

- Online regression

## ◆ Perceptrons

- Online SVM
  - Large margin / hinge loss
- Has nice mistake bounds (for separable case): see wiki
- In practice, use averaged perceptrons
- Passive Aggressive perceptrons and MIRA
  - Change  $w$  just enough to set its hinge loss to zero.